

Homework 6

GONG Kuiyuan

Student ID: 39-246182

Student name: GONG Kuiyuan

Preferred name: Eddie

Answers:

Task 1: I first set the working dictionary and use the function `read.csv` to access and store the data.

```
#setwd("~/Desktop/R for empirical research/HW/")
netflix <- read.csv("netflix_titles.csv")
```

Task 2: I use the verb `select()` together with `contains()` to select the columns I need. Here I selected columns that contain *at*. I also demonstrated the other approach where we can use `where(is.character)` to filter those empty columns. I need to use `head()` to restrict the list shown, otherwise it can't not be rendered.

```
library(pacman)
p_load(tidyverse)
netflix |>
  select(contains("at")) |>
  head(5)
```

	date_added	rating	duration
1	September 25, 2021	PG-13	90 min
2	September 24, 2021	TV-MA	2 Seasons
3	September 24, 2021	TV-MA	1 Season
4	September 24, 2021	TV-MA	1 Season
5	September 24, 2021	TV-MA	2 Seasons

```
#netflix |> select(where(is.character))
```

Task 3: I choose to only print the new variables, otherwise the file can't be rendered due to the big tibble.

```
netflix_2 <- netflix |>
  mutate(
    length = if_else(type == "Movie", parse_number(duration), NA),
    seasons = if_else(type == "TV Show", parse_number(duration), NA)
  ) |>
  relocate(length, seasons, .before = 1)

netflix_2 |>
  select(length, seasons) |>
  head() |>
  print()
```


	length	seasons
1	90	NA
2	NA	2
3	NA	1
4	NA	1
5	NA	2
6	NA	1

```
#print(netflix_2)
```

Task 4: Yes, we can observe from the data that there are more TV shows and the number of seasons over the time. A large number of observations are clustered between the years 2000 and 2025. Hypothesis: the number of TV shows and their seasons proliferates as technologies advance over the time.

```
ggplot(
  data = netflix_2,
  mapping = aes(x = release_year, y = seasons)
) +
  geom_point()
```


Warning: Removed 6132 rows containing missing values or values outside the scale range (`geom_point()`).



Task 5: Yes, we can observe a similar trend with the above plot. The number of movies increase over the time, however, the length of movies is commonly from 0 to 150 minutes, mostly concentrate around 100 minutes. Hypothesis: the number of movies proliferate as technologies advance over the time, however, the length of movies doesn't change a lot but stabilize on a scale.

```
ggplot(
  data = netflix_2,
  mapping = aes(x = release_year, y = length)
) +
  geom_point()
```

Warning: Removed 2680 rows containing missing values or values outside the scale range (``geom_point()``).



Task 6: The data is not “tidy” because food, clothing and other categories are column names here. The “tidy” data should be values of a single “category” column, with corresponding amounts in another column. For example, in the column “category” there will be values like “food” and “clothing”.

```
library(pacman)
p_load(griffen, griffendata)
print(expenditure_data1)
```

```
# A tibble: 100 x 5
   id  food clothing housing alcohol
  <int> <dbl>   <dbl>   <dbl>   <dbl>
1     1    95     56    172      8
2     2   236    183     98     34
3     3   158    269    521     51
4     4    69     71    382     37
5     5    39    171    187     61
6     6   210    107    666      8
7     7   319     96    589     17
8     8   146    189    261     17
9     9    17    206    184     18
10    10    20    127    115     27
# i 90 more rows
```

Task 7:

```
expenditure_data1 |>
  mutate(total_expenditure = food + clothing + housing + alcohol) |>
  print()
```

A tibble: 100 x 6

	id	food	clothing	housing	alcohol	total_expenditure
	<int>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>
1	1	95	56	172	8	331
2	2	236	183	98	34	551
3	3	158	269	521	51	999
4	4	69	71	382	37	559
5	5	39	171	187	61	458
6	6	210	107	666	8	991
7	7	319	96	589	17	1021
8	8	146	189	261	17	613
9	9	17	206	184	18	425
10	10	20	127	115	27	289

i 90 more rows

Task 8:

```
expenditure_pivot <- expenditure_data1 |>
  pivot_longer(
    cols = contains(c("food", "clothing", "housing", "alcohol")),
    names_to = 'category',
    values_to = 'expenditure') |>
  print()
```

A tibble: 400 x 3

	id	category	expenditure
	<int>	<chr>	<dbl>
1	1	food	95
2	1	clothing	56
3	1	housing	172
4	1	alcohol	8
5	2	food	236
6	2	clothing	183
7	2	housing	98
8	2	alcohol	34

```

 9      3 food          158
10      3 clothing     269
# i 390 more rows

```

Task 9:

```

expenditure_pivot <- expenditure_pivot |>
  group_by(id) |>
  summarise(total_expenditure = sum(expenditure)) |>
  print()

```

```

# A tibble: 100 x 2
   id total_expenditure
  <int>          <dbl>
1     1          331
2     2          551
3     3          999
4     4          559
5     5          458
6     6          991
7     7         1021
8     8          613
9     9          425
10    10          289
# i 90 more rows

```

Task 10: I prefer to use `group_by()` and `summarise()` to get the target value because it only prints out the total value and mutate other information, which is clearer for me.

Task 11:

```

library(pacman)
p_load(griffen, griffendata)
print(expenditure_data2)

```

```

# A tibble: 100 x 201
   id item1 item2 item3 item4 item5 item6 item7 item8 item9 item10 item11
  <int> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>
1     1    73    15   153   279    22   171   126    90    72    17    73
2     2   607   116    41    67    37   242    13    14   113    56   312
3     3   139   335    37   177    84    69    73   337    63   102   197

```

```

4      4      66    612    215    200     60     43     87    693     36     60    113
5      5     132     12    165     34    178     96    124     67     29    270    115
6      6     221     59    141     75     36     88     13     57     54     20    970
7      7      55    286     94     53    215     89    104     24     54     73    104
8      8     106     21    137     49    240     36    130    244    102    179    187
9      9     115     45    190     81    231     57    157    208    106    160     48
10     10      29    256    419    156     64    291     31     50     75     89    102
# i 90 more rows
# i 189 more variables: item12 <dbl>, item13 <dbl>, item14 <dbl>, item15 <dbl>,
#   item16 <dbl>, item17 <dbl>, item18 <dbl>, item19 <dbl>, item20 <dbl>,
#   item21 <dbl>, item22 <dbl>, item23 <dbl>, item24 <dbl>, item25 <dbl>,
#   item26 <dbl>, item27 <dbl>, item28 <dbl>, item29 <dbl>, item30 <dbl>,
#   item31 <dbl>, item32 <dbl>, item33 <dbl>, item34 <dbl>, item35 <dbl>,
#   item36 <dbl>, item37 <dbl>, item38 <dbl>, item39 <dbl>, item40 <dbl>, ...

```

Task 12: There are many ways to know the number of columns of a tibble. For example, we can know the size of the tibble when we print it out, or, we can use `ncol()` to report the number for us.

```
ncol(expenditure_data2)
```

```
[1] 201
```

Task 13: Without pivoting it longer, we need to manually add every term inside of the function in order to get the total expenditures.

```
#expenditure_data2 |>
#mutate(
#total_expenditure = item1 + item2 + item3 + item4...+ item200) |>
#print()
```

Task 14:

```
expenditure_data2 |>
  pivot_longer(
    cols = contains("item"),
    names_to = 'category',
    values_to = 'expenditure') |>
  group_by(id) |>
  summarise(total_expenditure = sum(expenditure)) |>
  print()
```

```
# A tibble: 100 x 2
      id total_expenditure
  <int>      <dbl>
1     1      32210
2     2      31455
3     3      32024
4     4      33438
5     5      32593
6     6      31102
7     7      31411
8     8      31497
9     9      33735
10    10      32288
# i 90 more rows
```

Task 15: Because the data now is “tidy” in the sense that we can manipulate the data easily using `group_by(id)` and continue to do other operations according to our needs. It transforms multiple expenditure categories (like food...) from separate columns into rows under a single column.

Task 16:

```
full_cps |>
  filter(state == "Kentacky") |>
  group_by(year, education_category, race) |>
  summarise(
    n = n(),
    wage = mean(wage, na.rm = TRUE)) |>
  filter(n > 10)
```

```
Error in `group_by()` :
! Must group by variables found in `.data`.
x Column `race` is not found.
```

Task 17: First, there is not such variable named “race” in the data, which makes `group_by()` invalid. Second, there is a misspelling of “Kentacky” instead of “Kentucky”, which makes the `filter()` invalid. Third, the previous writing style is very difficult for other people to read and debug the code, it may potentially increase the probability of making mistakes.

```
names(full_cps)
```



```

[1] "age"          "year"          "wage"
[4] "hours_lastweek" "employed"      "education_category"
[7] "educ_years"    "black"         "white"
[10] "female"        "married"       "single"
[13] "divorced"      "state"         "region"
[16] "sampling_weight"

```

```

full_cps |>
  filter(state == "Kentucky")

```

```

# A tibble: 0 x 16
# i 16 variables: age <dbl>, year <int>, wage <dbl>, hours_lastweek <dbl>,
#   employed <dbl>, education_category <chr>, educ_years <dbl>, black <dbl>,
#   white <dbl>, female <dbl>, married <dbl>, single <dbl>, divorced <dbl>,
#   state <chr>, region <chr>, sampling_weight <dbl>

```

```

full_cps |>
  filter(state == "Kentucky")

```

```

# A tibble: 9,159 x 16
   age year wage hours_lastweek employed education_category educ_years black
  <dbl> <int> <dbl>         <dbl>    <dbl> <chr>          <dbl> <dbl>
1    53  2000 31.2             45      1 college         18     0
2    35  2010 17.9             25      1 college         18     0
3    28  1984 NA              NA      0 highschool      NA     0
4    46  2005  9.86            38      1 somecollege     14     1
5    27  1980  9.03            50      1 highschool     10     0
6    28  1989 12.1            74      1 highschool     12     0
7    36  1998 NA              NA      0 highschool      9     0
8    57  2008 NA              NA      0 somecollege     14     0
9    62  2014 NA              20      1 college         16     0
10   49  1982 25.6            36      1 highschool     12     0
# i 9,149 more rows
# i 8 more variables: white <dbl>, female <dbl>, married <dbl>, single <dbl>,
#   divorced <dbl>, state <chr>, region <chr>, sampling_weight <dbl>

```